

IMAT3104

Database Management and Programming

Relational database transactions,
concurrency and recovery

Objectives

- Explain the concept of a relational database transaction.
- Explain the problems of transaction concurrency
- Examine the recovery of lost transactions

Introduction

- In database design we consider the support of transactions on the database.
- RDBMS facilities include:
 - Concurrency control.
 - Recovery.
 - Security.
 - Integrity.
- The transaction is central to understanding how a relational database is kept in a complete & consistent state.

Transactions

- Fundamental purpose of a database system is to carry out transactions.
- A transaction is a logical unit of work.
- A transaction consists of the execution of an application-specified sequence of operations (e.g SELECT, INSERT, UPDATE, DELETE).

Transaction Example

- One transaction with two operations:

```
UPDATE bank_account  
SET Balance = Balance - 100  
WHERE AccountNo = 348973;
```

```
UPDATE bank_account  
SET Balance = Balance + 100  
WHERE AccountNo = 678453;
```

- If the transaction failed between these operations, the database would be left in an inconsistent state.

Transaction ACID Properties

- **Atomicity** : An “All or Nothing” unit either executed in its entirety or not at all.
- **Consistency** : Transforms the database from one consistent state into another consistent state.
- **Isolation** : Executes independently of other transactions.
- **Durability** : The effects of a successfully completed transaction are permanently recorded in the database and must not be lost.

COMMIT/ROLLBACK protocol

- A transaction has one of two outcomes:
 - It completes successfully and data is committed to the database,
or
 - It is unsuccessful and data is rolled back to restore the database to its original consistent state.
- The user must define the boundaries of a transaction:
 - Begin the transaction : **START TRANSACTION**
 - Make changes to database permanent : **COMMIT**
or
 - Undo effects of transaction : **ROLLBACK**

Transaction Example

```
CREATE PROCEDURE bank_payment_transfer (intAccNoFrom
    INT, intAccNoTo INT, numPayment NUMERIC(10,2))
BEGIN
    DECLARE EXIT HANDLER FOR SQLEXCEPTION ROLLBACK;
    START TRANSACTION;
        UPDATE bank_account
            SET dbBalance = dbBalance - numPayment
            WHERE dbAccNo = intAccNoFrom;
        UPDATE bank_account
            SET dbBalance = dbBalance + numPayment
            WHERE dbAccNo = intAccNoTo;
    COMMIT;
END;
```


Log files

- Log files keep track of database transactions.
- Contain information about all changes to a database.
- Support the COMMIT/ROLLBACK protocol
- Log files contain:
 - **Before-image** of a data item, i.e. value before change.
 - **After-image** of a data item, i.e. value after change.
 - Lots of other information ...
- ROLLBACK uses before-image values backwards.
- COMMIT uses after-image values forwards.

Concurrency Control

- The process of managing simultaneous transactions on the database without having them interfere with one another.
- Concurrency allows much more efficient use of resources, resulting in improved database performance.
- However, concurrent access can cause interference and result in data inconsistencies

Concurrency Problems

- Three classic types of potential problems caused by concurrency:
 - **Lost update**
 - **Uncommitted dependency**
 - **Inconsistent analysis**

Lost Update problem

Time	Transaction 1	Transaction 2	Acc
t1	SELECT Acc		£100
t2		SELECT Acc	£100
t3	UPDATE Acc+£100		£200
t4		UPDATE Acc-£100	£0

- What happens to Transaction 1's update?

Uncommitted Dependency problem

Time	Transaction 3	Transaction 4	Acc
t1		UPDATE Acc+£200	£150
t2	SELECT Acc		£150
t3		ROLLBACK	-£50

- Transaction 3 is dependent on an uncommitted change at time t2.
- What is the consequence?

Inconsistent Analysis problem

Three accounts with three amounts (£):

1. ACC_1 400 2. ACC_2 500 3. ACC_3 300 sum = 1200

Time	Transaction 5	Transaction 6
t1	SELECT ACC_1 sum = 400	
t2	SELECT ACC_2 sum = 900	
t3		SELECT ACC_3
t4		UPDATE ACC_3 300 >> 200
t5		SELECT ACC_1
t6		UPDATE ACC_1 400 >> 500
t7		COMMIT
t8	SELECT ACC_3 sum = 1100 !!	

Serialisation

- If transactions were only allowed to execute in **series** then these concurrency problems would not exist.
- The **Serialisation** principle:
When two transactions operate in a parallel interleaved manner, then their effects should be the same as if they operated in a purely serial manner.
- We can guarantee serialisability by using the **two-phase locking** protocol ...

Locks

- A transaction's **lock** on data can prevent access to other transactions to prevent incorrect results.
- Two types of lock:
 - **Shared read lock** (S-lock)
 - shared reading but not writing.
 - **Exclusive write lock** (X-lock)
 - exclusive access.
- A transaction may have to **wait** for a lock to be released until placing its own lock on the data item.
- COMMIT & ROLLBACK release locks.

Two Phase Locking

- All locking operations must precede the first unlock operation in the transaction.
- There are two phases:
 - **Growing phase** - acquires all locks, but releases none.
 - **Shrinking phase** - releases all locks, but acquires none.

Lost Update solution

Time	Transaction 1	Transaction 2	Acc
t1	SELECT Acc (S-lock)		£100
t2		SELECT Acc (S-lock)	£100
t3	UPDATE Acc+100 (X-lock ?)		
t4	wait	UPDATE Acc-100 (X-lock ?)	
t5	wait	wait	

- No lost update, but a deadlock (more later)

Uncommitted Dependency solution

Time	Transaction 3	Transaction 4	Acc
t1		UPDATE Acc+200 (X-lock)	£150
t2	SELECT Acc (S-lock?)		
t3	wait	ROLLBACK (release X-lock)	-£50
t4	resume SELECT Acc (S-lock)		-£50

- Transaction 3 reads correct value of Account

Inconsistent Analysis solution

Three accounts with three amounts:

1. ACC_1= 400 2. ACC_2= 500 3. ACC_3=300 sum = 1200

Time	Transaction 5	Transaction 6
t1	SELECT ACC_1 (S-lock)	
	sum = 400	
t2	SELECT ACC_2 (S-lock)	
	sum = 900	
t3		SELECT ACC_3 (S-lock)
t4		UPDATE ACC_3 (X-lock)
		300 >> 200
t5		SELECT ACC_1 (S-lock)
t6		UPDATE ACC_1 (X-lock?)
t7		wait wait
t8	SELECT ACC_3 (S-lock?)	
t9	wait	wait

Deadlock

- When two or more transactions are in a simultaneous wait state, each waiting for locks held by the other to be released.
- Breaking the deadlock involves the “lock manager” choosing one of the transactions & rolling it back – this releases its locks and allows the other transaction to proceed.
- Some systems will then re-execute the rolled back transaction.

Starvation

- When a transaction waits for an indefinite period of time while others continue normally.
- For example, a transaction wants an X-lock which it can never get because of new transactions continually gaining S-locks
- Solutions include locking in the order requested or increasing priority according to length of waiting time and number of previous aborts.

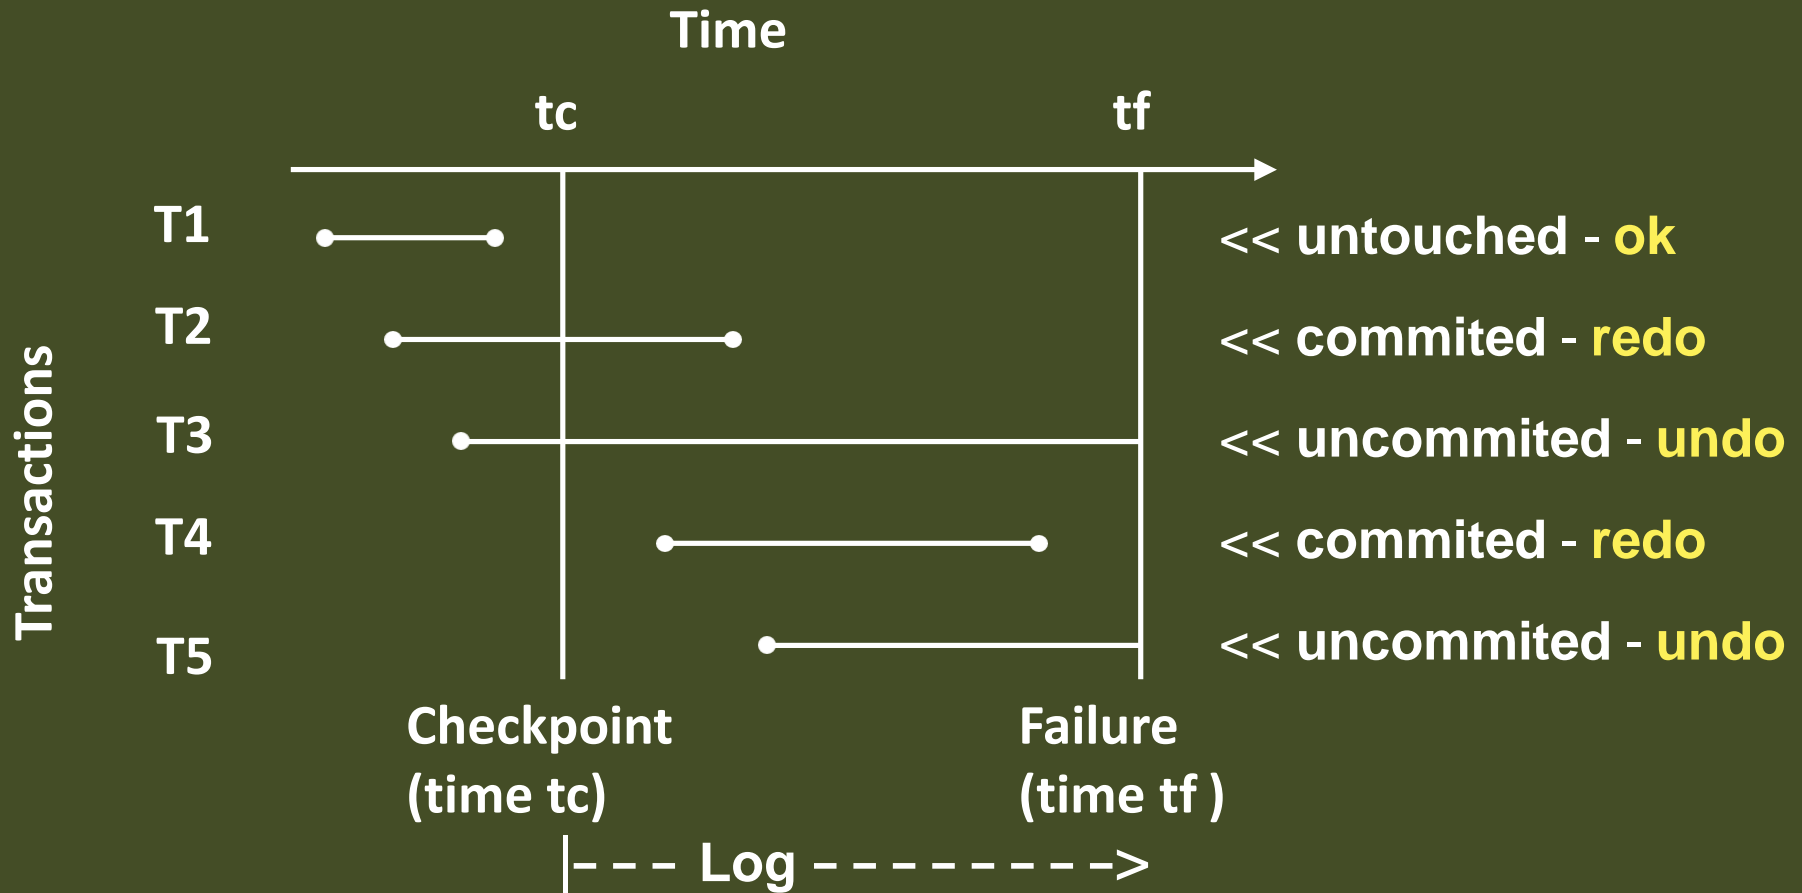
Recovery

- Restoring the database to a correct state in the event of failure.
- The transaction represents the basic unit of recovery.
- Recovery must ensure *atomicity* by undoing uncommitted transactions and *durability* by redoing committed transactions after failure.
- Recovery makes use of log files to restore the database.

Checkpoint

- A synchronisation point between the database and log file. All buffers are forced-written to secondary storage.
- Involves the following operations:
 - Write all log records to secondary storage.
 - Write contents of buffers to secondary storage.
 - Write a checkpoint record to the log file. This includes a list of all transactions in progress.

Using Checkpoints



Recovery technique

- Write ahead to the log file before the database as follows:
 - When transaction starts, register in log.
 - When write operation performed, write data record to log.
 - Once the log record is written, write the update to the database buffers.
 - The updates to the database itself are written when the buffers are next flushed to secondary storage.
 - When the transaction commits, register in log.
- Recovery uses the log to undo and redo transactions.

Summary

- Transactions and transaction control.
 - function, properties, COMMIT/ROLLBACK protocol, log file.
- The problems of transaction concurrency and how they can be controlled.
 - three common problems, two phase locking & deadlock
- The recovery of lost transactions.
 - checkpoint & write-ahead log protocol

Further Reading

- ROLLAND, F.D. (1998) *The Essence of Databases*, London, Prentice-Hall, Chapter 9.
- DATE, C.J. (2004) *An Introduction to Database Systems*, 8th ed., Reading, Addison-Wesley, Chapters 15 & 16.
- CONNOLLY, T.M. & BEGG, C.E. (2010) *Database Systems*, 5th. ed., Reading, Addison-Wesley, Chapter 20.